



# Semantic Analysis of Kaggle Dataset for Amazon Reviews

## CSCI 561 - Homework 2

*Shamim Samadi*

**Goal:** perform semantic analysis on Amazon reviews in Python

- Determination of whether a review is “**positive**” or “**negative**”
- binary classification problem

## Steps:

### 1. Pre-processing:


- Bag-of-words representation
- tokenization, remove punctuation, remove the stop words, lower case
- train-test split: every 5<sup>th</sup> sample belongs to test, the rest is training data

### 2. Classifier training

- Naive Bayes
- Decision Tree
- Neural Networks (multilayer perceptron)

### 3. Classifier evaluation on test data

- F-1 score, precision, recall
- ROC curve, AUC score



Naive Bayes was chosen as the 3<sup>rd</sup> classifier because:

- model is simple and computationally efficient
- a popular choice for text classification applications
- the feature independence assumption seems to go well with the bag-of-words representation

# Results & Discussion:



Precision

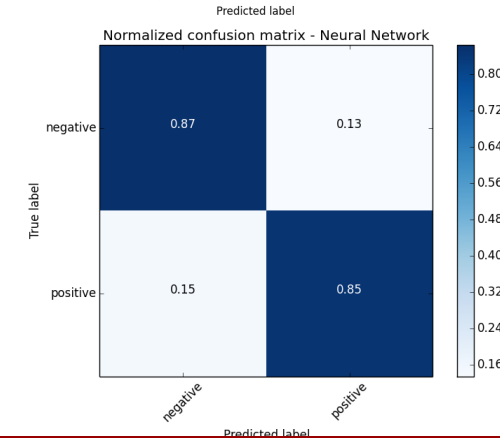
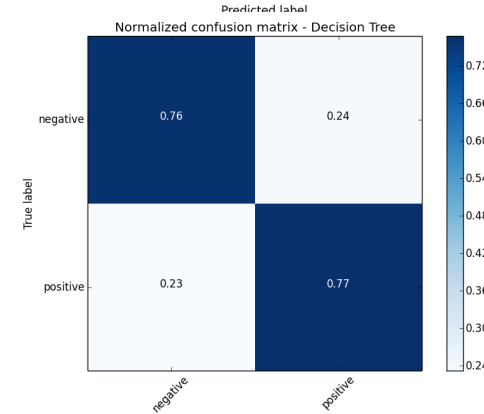

	class 0 (negative)	class 1 (positive)	average/total
<b>Naive Bayes</b>	0.84	<b>0.86</b>	0.85
<b>Decision Tree</b>	0.76	0.77	0.77
<b>Neural Network</b>	<b>0.85</b>	<b>0.86</b>	<b>0.86</b>

Recall

	class 0 (negative)	class 1 (positive)	average/total
<b>Naive Bayes</b>	0.86	0.84	0.85
<b>Decision Tree</b>	0.76	0.77	0.77
<b>Neural Network</b>	<b>0.87</b>	<b>0.85</b>	<b>0.86</b>

F1-Score

	class 0 (negative)	class 1 (positive)	average/total
<b>Naive Bayes</b>	0.85	0.85	0.85
<b>Decision Tree</b>	0.76	0.77	0.77
<b>Neural Network</b>	<b>0.86</b>	<b>0.86</b>	<b>0.86</b>




# Results & Discussion:

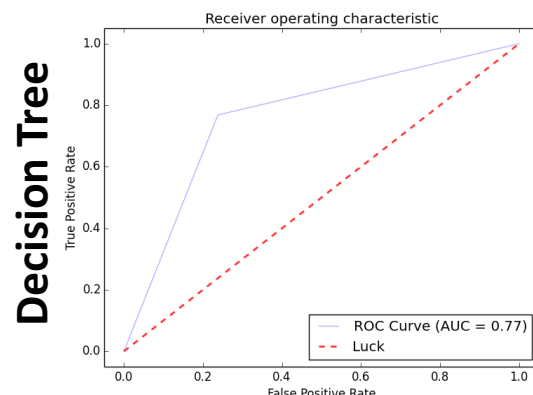
- Performance ranking
  1. Neural Network
  2. Naive Bayes
  3. Decision Tree
- Although neural network classifier has a slight edge in performance when it comes to evaluation metrics, the Naive Bayes classifier would be an overall better choice since neural network requires intensive training and is computationally very expensive.
- Decision tree performs poorly in high-dimensional feature spaces (which is the case in the bag-of-words feature representation used in this homework)
  - tricky to find the “key” tokens (as the key nodes) to split the tree on



**AUC  
= 0.92**



**AUC  
= 0.77**



**AUC  
= 0.93**

